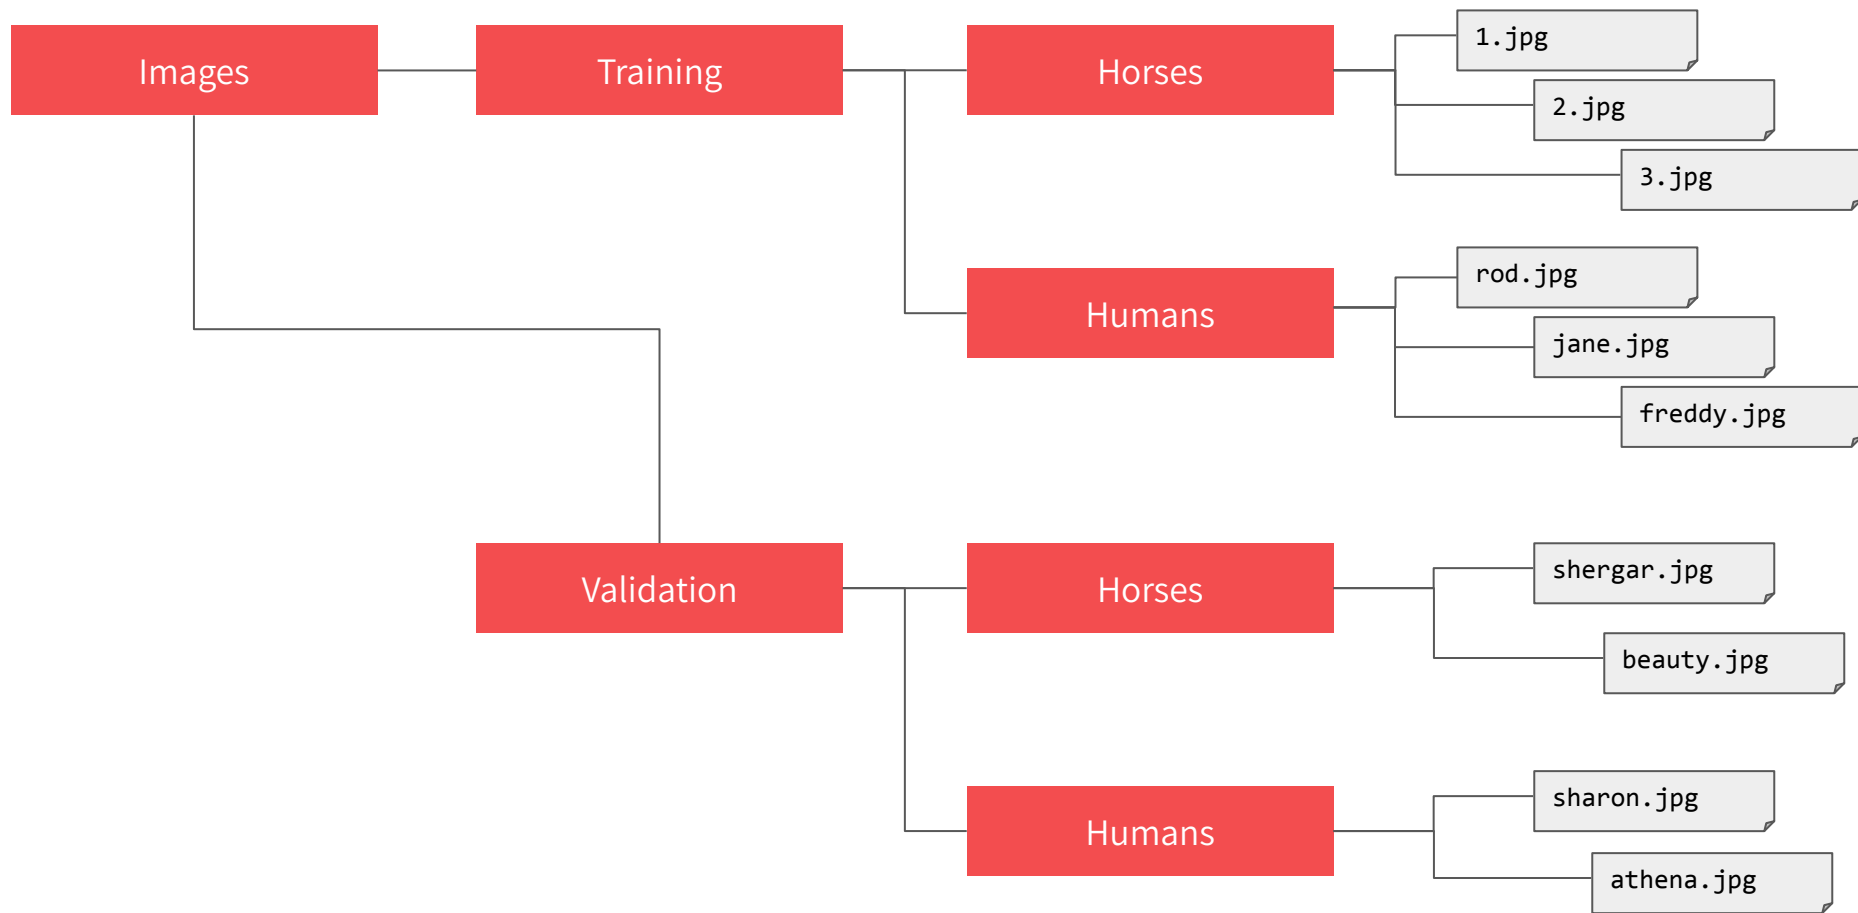


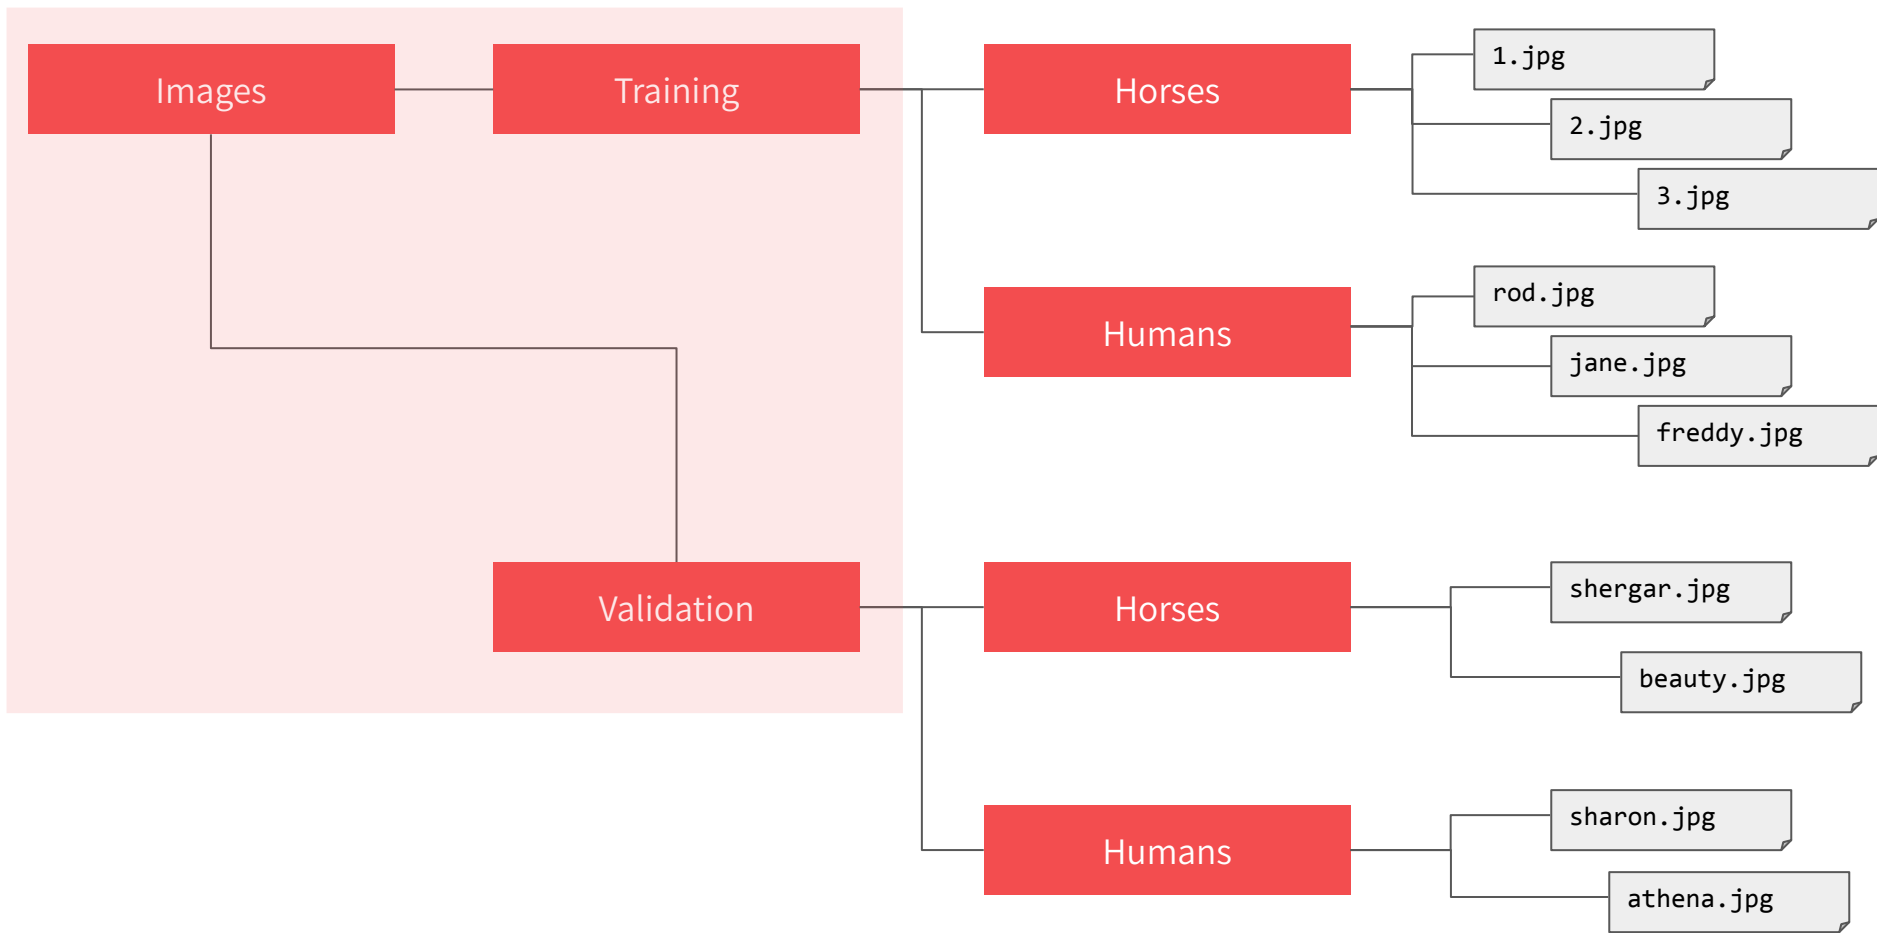
# tf.data API

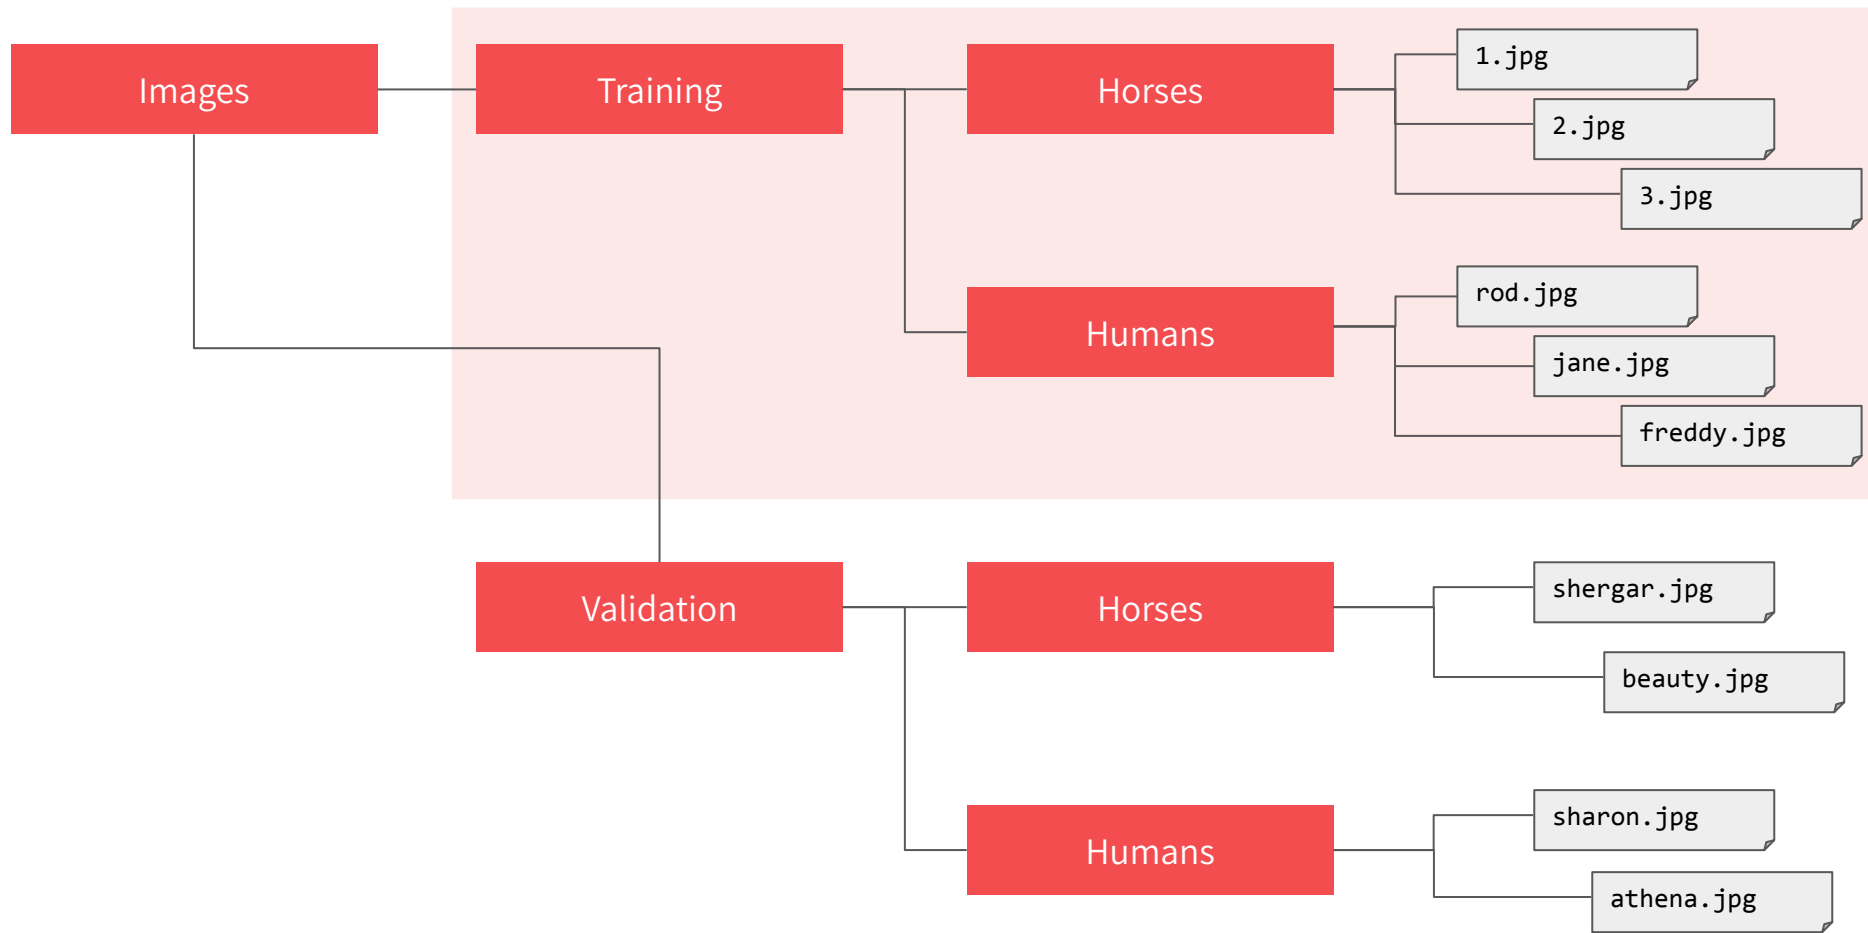


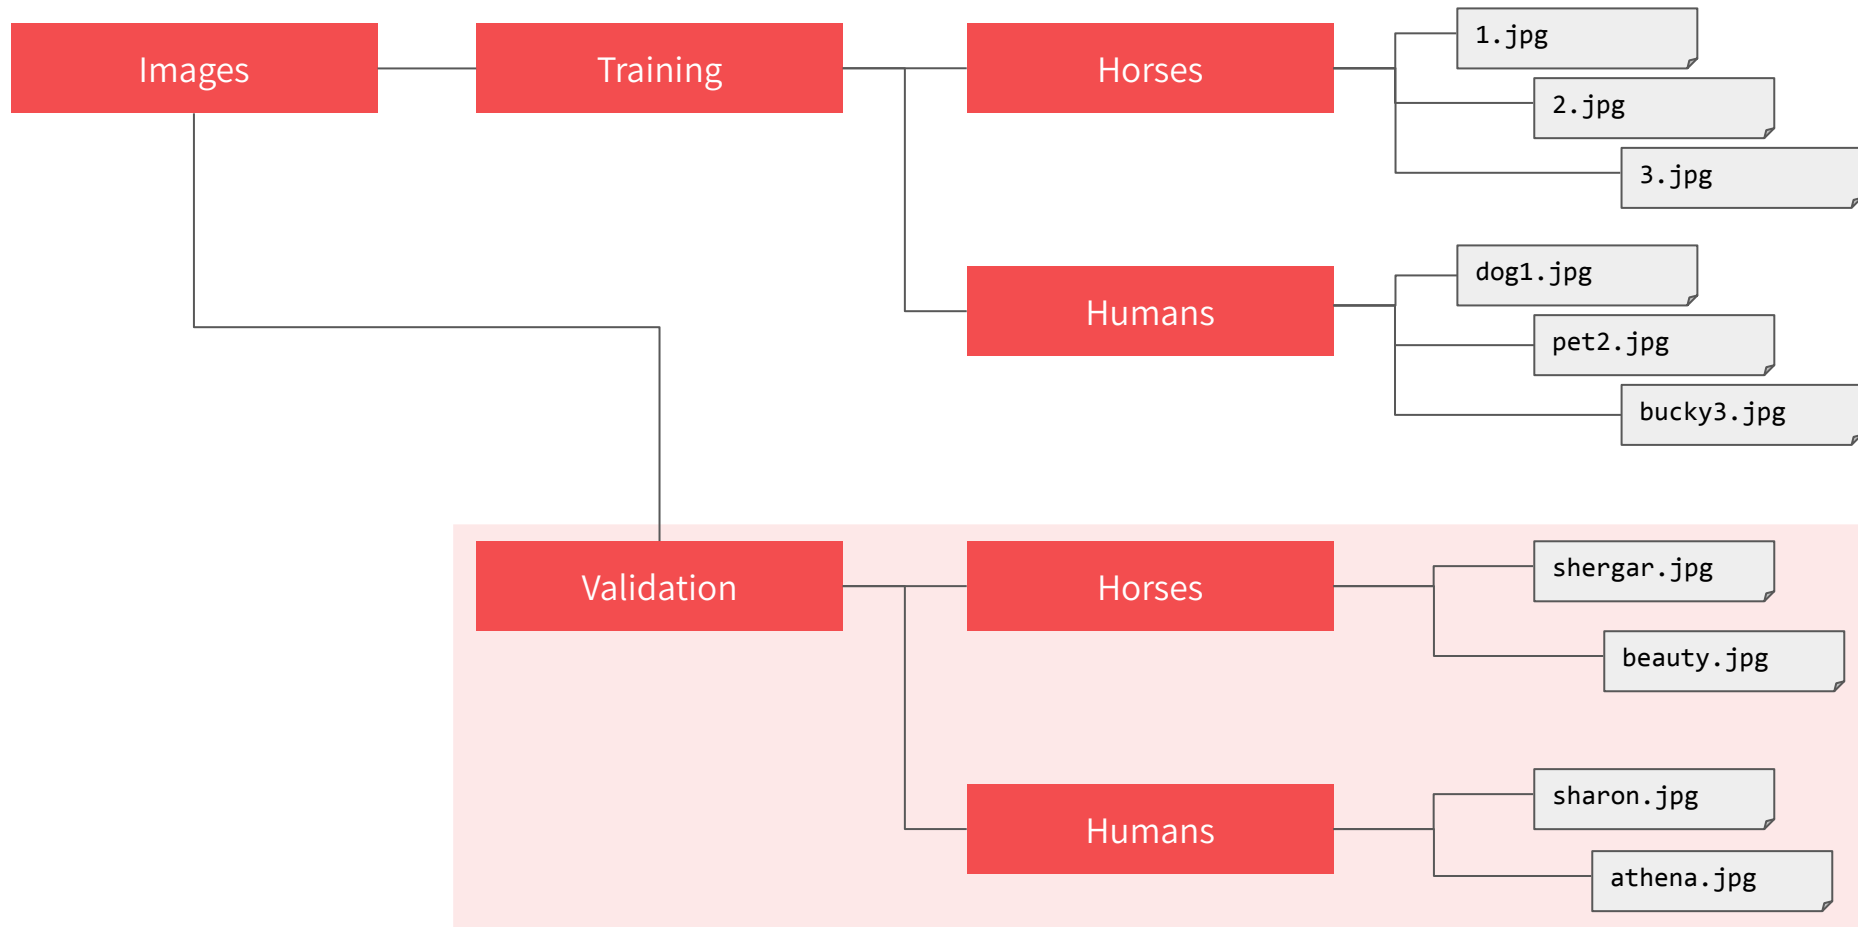
**tf.keras.utils.image\_dataset\_from\_directory**











```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```





```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



```
train_dataset = tf.keras.utils.image_dataset_from_directory(  
    TRAIN_DIR,  
    image_size=(300, 300),  
    batch_size=128,  
    label_mode='binary'  
)
```



# **tf.data.Dataset**

[https://www.tensorflow.org/api\\_docs/python/tf/data/Dataset](https://www.tensorflow.org/api_docs/python/tf/data/Dataset)

```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)
```

```
train_dataset_scaled = train_dataset.map(  
    lambda image, label: (rescale_layer(image), label))
```



```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)

train_dataset_scaled = train_dataset.map(
    lambda image, label: (rescale_layer(image), label))
```



```
rescale_layer = tf.keras.layers.Rescaling(scale=1./255)

train_dataset_scaled = train_dataset.map(
    lambda image, label: (rescale_layer(image), label))
```



```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```

```
train_dataset_final = (train_dataset_scaled  
                        .cache()  
                        .shuffle(buffer_size=1000)  
                        .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    VAL_DIR,
    image_size=(300, 300),
    batch_size=32,
    label_mode='binary'
)

validation_dataset_scaled = validation_dataset.map(lambda image, label:
(rescale_layer(image), label))

# Configure the validation dataset
validation_dataset_final = (validation_dataset_scaled
    .cache()
    .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
validation_dataset = tf.keras.utils.image_dataset_from_directory(
    VAL_DIR,
    image_size=(300, 300),
    batch_size=128,
    label_mode='binary'
)
```

```
validation_dataset_scaled = validation_dataset.map(lambda image, label:
    (rescale_layer(image), label))
```

```
# Configure the validation dataset
```

```
validation_dataset_final = (validation_dataset_scaled
    .cache()
    .prefetch(buffer_size=tf.data.AUTOTUNE))
```



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Input(shape=(300, 300, 3)),  
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(512, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```



```
model = tf.keras.models.Sequential([
    tf.keras.layers.Input(shape=(300, 300, 3)),
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d_5 (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_6 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_6 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_7 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_7 (MaxPooling2D)	(None, 35, 35, 64)	0
flatten_1 (Flatten)	(None, 78400)	0
dense_2 (Dense)	(None, 512)	40141312
dense_3 (Dense)	(None, 1)	513
Total params: 40,165,409		
Trainable params: 40,165,409		
Non-trainable params: 0		

```
model.compile(loss='binary_crossentropy',  
              optimizer=tf.keras.optimizers.RMSprop(learning_rate=0.001),  
              metrics=['accuracy'])
```

<https://youtu.be/zLRB4oupj6g>



Machine Learning



0:06 / 8:58



2.1.4 Gradient Descent in Practice II Learning Rate by Andrew Ng

```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```





```
history = model.fit(  
    train_dataset_final,  
    epochs=15,  
    validation_data=validation_dataset_final,  
    verbose=2)
```



```
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():

    # predicting images
    path = '/content/' + filename
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
    image = tf.keras.utils.img_to_array(image)
    image = rescale_layer(image)
    image = np.expand_dims(image, axis=0)

    prediction = model.predict(image, verbose=0)[0][0]
    print(f'\nmodel output: {prediction}')

    if prediction > 0.5:
        print(filename + " is a human")
    else:
        print(filename + " is a horse")
```



```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for filename in uploaded.keys():
```

```
    # predicting images
```

```
    path = '/content/' + filename
```

```
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
```

```
    image = tf.keras.utils.img_to_array(image)
```

```
    image = rescale_layer(image)
```

```
    image = np.expand_dims(image, axis=0)
```

```
    prediction = model.predict(image, verbose=0)[0][0]
```

```
    print(f'\nmodel output: {prediction}')
```

```
    if prediction > 0.5:
```

```
        print(filename + " is a human")
```

```
    else:
```

```
        print(filename + " is a horse")
```



```
from google.colab import files
```

```
uploaded = files.upload()
```

```
for filename in uploaded.keys():
```

```
    # predicting images
```

```
    path = '/content/' + filename
```

```
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
```

```
    image = tf.keras.utils.img_to_array(image)
```

```
    image = rescale_layer(image)
```

```
    image = np.expand_dims(image, axis=0)
```

```
    prediction = model.predict(image, verbose=0)[0][0]
```

```
    print(f'\nmodel output: {prediction}')
```

```
    if prediction > 0.5:
```

```
        print(filename + " is a human")
```

```
    else:
```

```
        print(filename + " is a horse")
```



```
from google.colab import files

uploaded = files.upload()

for filename in uploaded.keys():

    # predicting images
    path = '/content/' + filename
    image = tf.keras.utils.load_img(path, target_size=(300, 300))
    image = tf.keras.utils.img_to_array(image)
    image = rescale_layer(image)
    image = np.expand_dims(image, axis=0)

    prediction = model.predict(image, verbose=0)[0][0]
    print(f'\nmodel output: {prediction}')

    if prediction > 0.5:
        print(filename + " is a human")
    else:
        print(filename + " is a horse")
```

